

Introduction to Batch Scheduling

ICS632: Principles of High-Performance Computing

Henri Casanova (henric@hawaii.edu)


Fall 2015

Foreword

- In these lecture notes we talk about batch scheduling
- This is something we are confronted to when using most HPC platforms
- We'll review basic principles, algorithms, and results
- We'll make small connections to theory

Batch Scheduling

- Many users typically use the same cluster
- They shouldn't step of each others' toes
 - They should use *dedicated* subsets of the platform
- A **Batch Scheduler** is responsible for the sharing



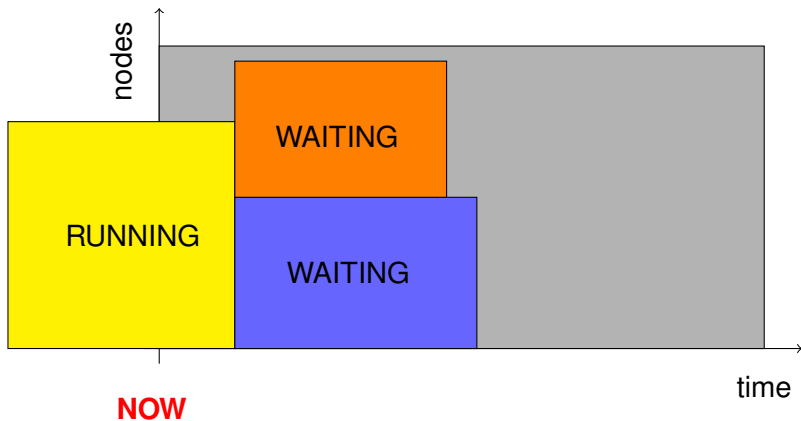
Batch Scheduler

- The batch scheduler keeps track of:
 - Currently running jobs
 - Jobs in the queue
- A job is defined by:
 - Arrival time
 - Requested number of nodes and duration (“4 nodes for 2 hours”)
- Tons of bells and whistles
 - Multiple queues for various classes of jobs
 - Advance reservations
 - Charging policies
 - Heterogeneous nodes
 - ...
- Let's keep it simple

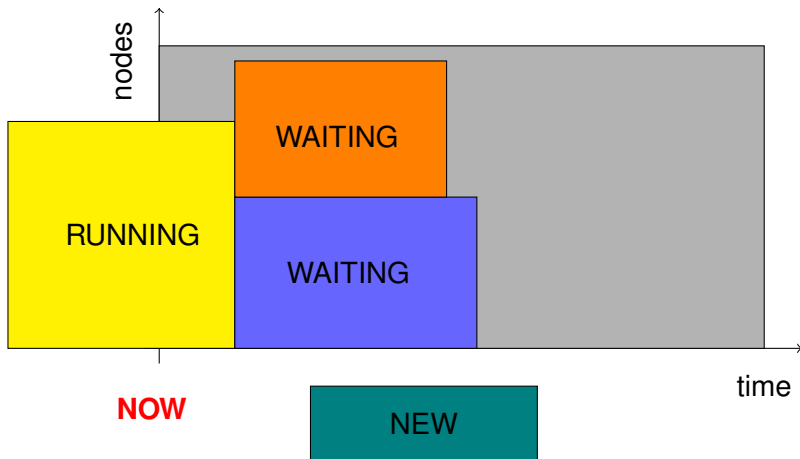
Graphical Representation of a Schedule



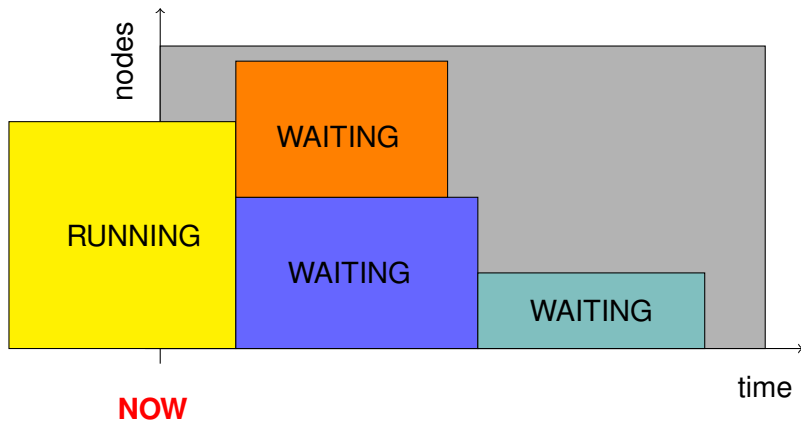
Graphical Representation of a Schedule



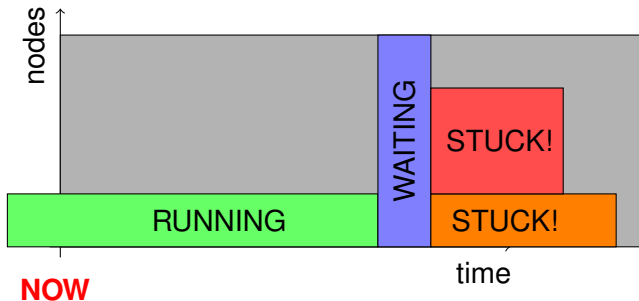
Graphical Representation of a Schedule



Graphical Representation of a Schedule

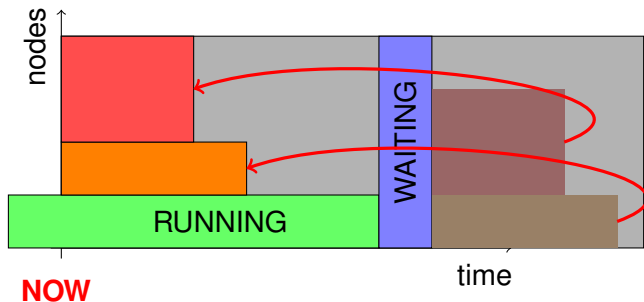


First-Come-First-Serve (FCFS)



- Problem: fragmentation leads to idle time
 - Not good for users, not good for cluster owners

Solution: Backfilling



- Allow jobs to “jump in line”
 - Sometimes happens at the supermarket checkout!

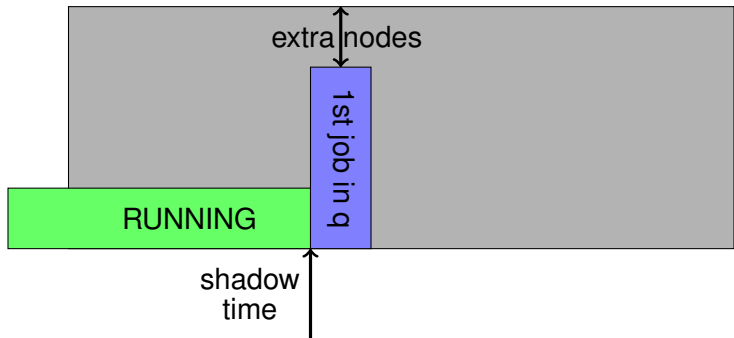
Backfilling?

- Question: which jobs should be picked for promotion through the queue?
 - Clearly jobs that are short/small
 - Hence the motivation to provide not-too-conservative estimates through the `-t` option in SLURM
- Many heuristics are possible, but two have been studied in detail:
 - EASY
 - Conservative Back Filling (CBF)
- Most production batch schedulers are EASY-ish
 - A few systems use CBF (OAR)


EASY Backfilling

- Extensible Argonne Scheduling System
- Maintain only one "reservation," for the first job in the queue
- Definitions:
 - Shadow time: time at which the first job in the queue starts execution
 - Extra nodes: number of nodes idle when the first job in the queue starts execution
- Go through the queue in order starting with the 2nd job
- Backfill a job if:
 - It will terminate by the shadow time, or
 - it needs less than the extra nodes


EASY Example



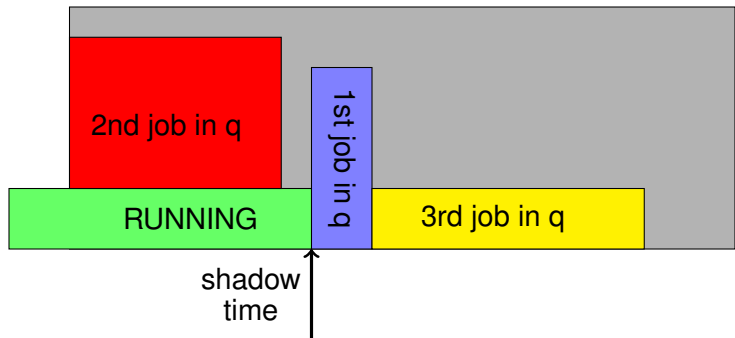
EASY Example



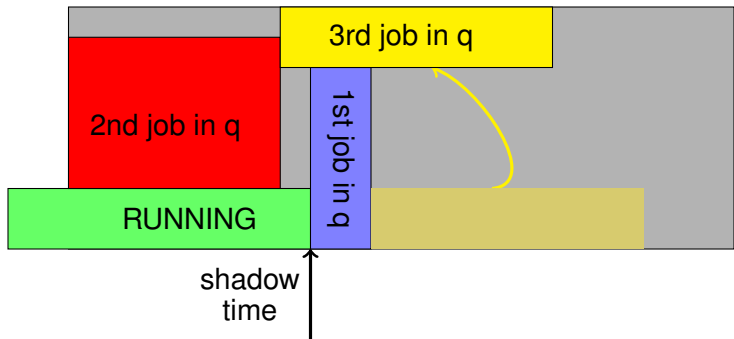
EASY Example



EASY Example




EASY Example




EASY Properties

- Unbounded delay
 - The first job in the queue will never be delayed by backfilled jobs
 - BUT other jobs may be delayed infinitely!
- Let's see how that can happen...


EASY: Unbounded Delay



EASY: Unbounded Delay



EASY: Unbounded Delay



EASY Properties

- Unbounded delay
 - The first job in the queue will never be delayed by backfilled jobs
 - BUT other jobs may be delayed infinitely!
- But there is no starvation:
 - Delay of first job is bounded by runtime of current jobs
 - When the first job runs, the second job becomes the first job in the queue
 - Once it is the first job, it cannot be delayed further

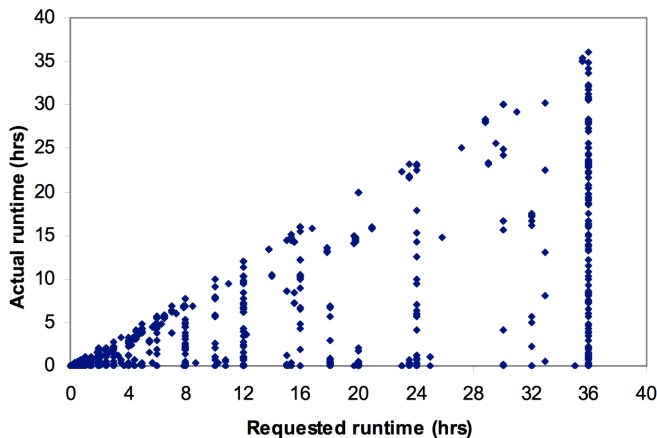
Conservative Backfilling

- EVERY job has a "reservation"
- A job may be backfilled only if it does not delay any other job ahead of it in the queue
- Fixes the unbounded delay problem that EASY has
 - EASY favors small long jobs
 - EASY harms large short jobs
- More complicated to implement
 - The algorithm must find holes in the schedule

Backfilling Decisions

- When does backfilling happen?
 - When a new job arrives
 - When the first job in the queue starts
 - When a job finishes (early!)
- All scheduling decisions are based on job durations
- Job durations are provided by users
- A job is killed if it goes over
- Trade-off:
 - Provide an aggressive estimate: you may be killed!
 - Provide a conservative estimate: you won't be backfilled easily
- Question: Are job duration estimates accurate?

Estimates are NOT Accurate



from *Are user runtime estimates inherently inaccurate?*

Bailey Lee et al., JSSPP'04

Batch Scheduling Research

- Most parallel platforms are managed by batch schedulers or things like batch schedulers
- There are many difficult questions
 - e.g., how do you schedule without knowing job durations?
- The people managing the platform often want to maximize resource utilization/profit
- The users want to minimize response time
- Many “tricks” can be played:
 - Picking the right “shape” so that you’ll be backfilled
 - Chop up your job into multiple pieces
 - Aggressively submit versions of the same job (different shapes), perhaps to multiple systems, and cancel when one begins
 - ...
- What do theoreticians think?

Whats a Good Batch Schedule?

- We have an **on-line scheduling problem**
- The first step is to define a metric of goodness of a schedule
- Let's look at a few likely metrics
- **Wait time:** time spent in the queue
 - Wait time is annoying, so likely a good thing to minimize
 - Not a great idea:
 - Job #1 needs 100h on 1000 nodes and waits 1h
 - Job #2 needs 1s on 1 node and waits 1h
 - Clearly Job #1 is really happy, and Job #2 is not happy at all

Whats a Good Batch Schedule?

- **Turn-around time:** Wait time + Execution time
 - Called *flow time* by theoreticians
 - Not a great idea:
 - Job #1 needs 1h of compute time and waits 1s
 - Job #2 needs 1s of compute time and waits 1h
 - Clearly Job #1 is really happy, and Job #2 is not happy at all

Whats a Good Batch Schedule?

- What we want is a metric that represents “happiness” for small, large, short, long jobs
- **Slowdown:** $(\text{Wait time} + \text{Execution time}) / \text{Execution time}$
 - Called *stretch* by theoreticians
 - Quantifies loss of performance due to competition for the processors
 - Takes care of the short vs. long job problem
 - Doesn't really say anything about job size...
- Two possible objectives:
 - minimize the *sum stretch* (make jobs happy on average)
 - minimize the *max stretch* (make the least happy job as happy as possible)

On-line Max Stretch: Difficult

- The offline scheduling problem is NP-complete
- On 1 processor, with preemption allowed, there is a $O(\sqrt{X})$ -competitive algorithm
 - X is the ratio of largest to smallest job duration
 - Competitive ratio: ratio to the performance of an off-line algorithm that knows all jobs
- Without preemption, no approximation algorithm exists

On-line Sum Stretch: “Easier”

- The offline scheduling problem is NP-complete
- The SRPT (Shortest Remaining Processing Time) with preemption is 2-competitive
 - At each instant, run the job that's the closest to finishing
 - Preemption overhead is zero
- Theorems that show that an algorithm good at minimizing sum stretch can be arbitrarily bad at minimizing max stretch, and conversely

Conclusion: A Massive Divide

- Practice: No preemption in batch scheduling, need for many scheduling configuration knobs
- Theory: Without preemption, we can't do anything guaranteed anyway
- The two remain very divorced
- The stretch is used as a metric to evaluate how good scheduling is in practice, but it often isn't the objective of the batch scheduler
 - That objective is complex, mysterious, and not necessarily theoretically-motivated